

# AucklandHousePrices

July 26, 2020

## 1 Auckland House Prices Analysis

**Executive Summary** The dataset is based in Auckland the objective to create a linear regression to predict house price based on features in the data, which was provided in the MSA phase 1 data-science pathway. It contains all a variety of information regarding houses in Auckland along with price evaluation and land and bedroom information as well as longitude and latitude information. However, some information related to the population for housing was missing. So API call to a 2018 census dataset was conducted using longitude, latitude through a vector query and SA1 from the initial dataset for each house and a new column was inserted into the dataset. Additionally, the depression index was also added to the primary dataset using Otago researched-based dataset on house depression in Auckland. Finally, some data clean was completed; Nan values were replaced with estimates. Some of the data in the housing data set contained categorical data which required transformations for plotting visualisations. These were suburbs and address.

The analysis is based on 981 observations for each of the 17 variables. A correlation matrix was used initially to see if any of the 17 variables were related price CV (estimate of the house price) if they had a strong relationship. Before this, descriptive statistics was used initially to determine the spread of each variable. Additionally, outliers were removed with an interquartile range metric using land Area. Three correlated values were found; these were Bedrooms, Bathrooms, NZDep2018. Finally, a linear regression model was then fitted to predict CV based on these inputs.

```
[309]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import json
import sys
import pandas as pd
import requests
import time
%matplotlib inline
```

```
[311]: #different suburbs effect prices bedroom size of house land etc givin a price
→, suburbs, rooms predict the price
# using regression good analysis techniques using to determine
house_prices = pd.read_csv('DatasetforAssignment.csv')
#remove data types i.e object to float
house_prices.head()
```

```
[311]: Bedrooms Bathrooms Address Land area \
0      5      3.0 106 Lawrence Crescent Hill Park, Auckland 714
1      5      3.0      8 Corsica Way Karaka, Auckland 564
2      6      4.0 243 Harbourside Drive Karaka, Auckland 626
3      2      1.0 2/30 Hardington Street Onehunga, Auckland 65
4      3      1.0 59 Israel Avenue Clover Park, Auckland 601
```

```
CV Latitude Longitude SA1 0-19 years 20-29 years \
0 960000 -37.012920 174.904069 7009770 48 27
1 1250000 -37.063672 174.922912 7009991 42 18
2 1250000 -37.063580 174.924044 7009991 42 18
3 740000 -36.912996 174.787425 7007871 42 6
4 630000 -36.979037 174.892612 7008902 93 27
```

```
30-39 years 40-49 years 50-59 years 60+ years Suburbs
0      24      21      24      21 Manurewa
1      12      21      15      30 Karaka
2      12      21      15      30 Karaka
3      21      21      12      15 Onehunga
4      33      30      21      33 Clover Park
```

```
[313]: #remove convert colcumn to float and remove chars in text
house_prices['Land area'] = house_prices['Land area'].str.extract('(\d+)').
→astype(float)
```

```
[315]: #load in the depresisation scores
dep_indexs = pd.read_excel('otago730395.xlsx')
dep_indexs.head()
```

```
[315]: SA12018_code NZDep2018 NZDep2018_Score URPopnSA1_2018 SA22018_code \
0      7000000      10.0      1245.0      141      100100
1      7000001      10.0      1245.0      114      100100
2      7000002      NaN      NaN      0      100300
3      7000003      10.0      1207.0      225      100100
4      7000004      9.0      1093.0      138      100100
```

```
SA22018_name
0      North Cape
1      North Cape
2 Inlets Far North District
3      North Cape
4      North Cape
```

```
[317]: #now we need to get make the requests from the API for population at Latitude
→and Longitude
def get_pop_at_lat_long(long,lat):
    url = 'https://koordinates.com/services/query/v1/vector.json'
    #from nz cenius data
    layer_id = 104612
```

```

params = {
    'key': '772a6ce37abd44a7a65fb8af6ad6ac28',
    'layer' : layer_id,
    'x':long,
    'y':lat,
    'format':'json'
}
response = requests.get(url,params = params)
if(response.status_code != 200):
    return response.status_code
return response.json()

# handles the reponse from the function above extracts relvatant information
def parse_response(input_json_response,sal):
    layer_id = '104612'
    cencius_data = _
    →input_json_response['vectorQuery']['layers'][layer_id]['features']
    for item in cencius_data:
        res_SA1 = item['properties']['SA12018_V1_00']
        if (int(res_SA1) == int(sal) ):
            C18_CURPop = item['properties']['C18_CURPop']
            return {"SA12018_V1_00":int(sal),"C18_CURPop":int(C18_CURPop)}
    return {"SA12018_V1_00":int(sal),"C18_CURPop":None}

# used to get population value
def extract_population(long,lat,sal):
    api_data_response = get_pop_at_lat_long(long,lat)
    population_value = parse_response(api_data_response,sal)
    return population_value["C18_CURPop"]

#extract the dep index from one dataframe ouput it SA1
def extract_depreciation(df_dep,sal):
    # get row where sal has the same value
    row = df_dep.loc[df_dep['SA12018_code'] == sal]
    return float(row['NZDep2018'])

```

```

[319]: sample_house_prices['C18_CURPop'] = sample_house_prices.apply(lambda row :_
    →extract_population(row['Longitude'],row['Latitude'],row['SA1']),axis = 1)
sample_house_prices.head()

```

```

[319]:
Bedrooms    Bathrooms                                Address    Land area \
0           5         3.0    106 Lawrence Crescent Hill Park, Auckland    714.0
1           5         3.0                                8 Corsica Way Karaka, Auckland    564.0
2           6         4.0        243 Harbourside Drive Karaka, Auckland    626.0
3           2         1.0    2/30 Hardington Street Onehunga, Auckland     65.0
4           3         1.0        59 Israel Avenue Clover Park, Auckland    601.0

CV    Latitude    Longitude    SA1    0-19 years    20-29 years \

```

0	960000	-37.012920	174.904069	7009770	48	27
1	1250000	-37.063672	174.922912	7009991	42	18
2	1250000	-37.063580	174.924044	7009991	42	18
3	740000	-36.912996	174.787425	7007871	42	6
4	630000	-36.979037	174.892612	7008902	93	27

	30-39 years	40-49 years	50-59 years	60+ years	Suburbs	C18_CURPop \
0	24	21	24	21	Manurewa	174
1	12	21	15	30	Karaka	129
2	12	21	15	30	Karaka	129
3	21	21	12	15	Onehunga	120
4	33	30	21	33	Clover Park	231

	NZDep2018
0	6.0
1	1.0
2	1.0
3	2.0
4	9.0

```
[321]: #adding the depreciation column
sample_house_prices['NZDep2018'] = sample_house_prices.apply(lambda row :
    →extract_depreciation(dep_indexes,row['SA1']),axis = 1)
sample_house_prices.head()
```

```
[321]: Bedrooms    Bathrooms    Address    Land area \
0          5         3.0  106 Lawrence Crescent Hill Park, Auckland    714.0
1          5         3.0           8 Corsica Way Karaka, Auckland    564.0
2          6         4.0  243 Harbourside Drive Karaka, Auckland    626.0
3          2         1.0  2/30 Hardington Street Onehunga, Auckland    65.0
4          3         1.0   59 Israel Avenue Clover Park, Auckland    601.0
```

	CV	Latitude	Longitude	SA1	0-19 years	20-29 years \
0	960000	-37.012920	174.904069	7009770	48	27
1	1250000	-37.063672	174.922912	7009991	42	18
2	1250000	-37.063580	174.924044	7009991	42	18
3	740000	-36.912996	174.787425	7007871	42	6
4	630000	-36.979037	174.892612	7008902	93	27

	30-39 years	40-49 years	50-59 years	60+ years	Suburbs	C18_CURPop \
0	24	21	24	21	Manurewa	174
1	12	21	15	30	Karaka	129
2	12	21	15	30	Karaka	129
3	21	21	12	15	Onehunga	120
4	33	30	21	33	Clover Park	231

	NZDep2018
0	6.0

```
1      1.0
2      1.0
3      2.0
4      9.0
```

Adding the two extra columns for houses price dataset

```
[323]: house_prices['C18_CURPop'] = house_prices.apply(lambda row :
    →extract_population(row['Longitude'],row['Latitude'],row['SA1']),axis = 1)
```

```
[339]: house_prices['NZDep2018'] = house_prices.apply(lambda row :
    →extract_depreciation(dep_indexs,row['SA1']),axis = 1)
```

```
[327]: house_prices.head()
```

```
[327]: Bedrooms  Bathrooms                Address  Land area \
0          5         3.0  106 Lawrence Crescent Hill Park, Auckland    714.0
1          5         3.0                8 Corsica Way Karaka, Auckland    564.0
2          6         4.0      243 Harbourside Drive Karaka, Auckland    626.0
3          2         1.0  2/30 Hardington Street Onehunga, Auckland     65.0
4          3         1.0    59 Israel Avenue Clover Park, Auckland    601.0
```

```
CV  Latitude  Longitude  SA1  0-19 years  20-29 years \
0  960000 -37.012920  174.904069  7009770         48         27
1 1250000 -37.063672  174.922912  7009991         42         18
2 1250000 -37.063580  174.924044  7009991         42         18
3  740000 -36.912996  174.787425  7007871         42          6
4  630000 -36.979037  174.892612  7008902         93         27
```

```
30-39 years  40-49 years  50-59 years  60+ years  Suburbs  C18_CURPop \
0          24          21          24          21  Manurewa         174
1          12          21          15          30   Karaka         129
2          12          21          15          30   Karaka         129
3          21          21          12          15  Onehunga         120
4          33          30          21          33  Clover Park         231
```

```
NZDep2018
0      6.0
1      1.0
2      1.0
3      2.0
4      9.0
```

The above data frame imported contains a variety of information for houses in Auckland. An investigation is being conducted to determine if we can estimate houses prices based on information related house features. To determine the relevant features, we are going to perform some visualisations with the data frame to look for patterns.

```
[328]: # check if any row has null value NEED to check for outliers
pd.isnull(house_prices[:]).any(axis=1)
nan_stuff = house_prices[house_prices.isnull().any(axis=1)]
```

```
print(nan_stuff)
```

	Bedrooms	Bathrooms	Address \
309	4	NaN	14 Hea Road Hobsonville, Auckland
311	4	NaN	16 Hea Road Hobsonville, Auckland
568	1	1.0	14 Te Rangitawhiri Road Great Barrier Island, ...

	Land area	CV	Latitude	Longitude	SA1	0-19 years \
309	214.0	1250000	-36.798371	174.647430	7002267	60
311	245.0	1100000	-36.798371	174.647430	7002267	60
568	2141.0	740000	-36.197282	175.416921	7001131	27

	20-29 years	30-39 years	40-49 years	50-59 years	60+ years \
309	66	60	24	24	18
311	66	60	24	24	18
568	6	6	18	39	60

	Suburbs	C18_CURPop	NZDep2018
309	Hobsonville	252	2.0
311	Hobsonville	252	2.0
568	NaN	156	9.0

[329]: *#from this we found NaN values in the dataset two bathroom entrys and Suburb*  
*→entry*  
 subs = house\_prices[house\_prices['Suburbs'] == 'Hobsonville']  
 print(subs)

	Bedrooms	Bathrooms	Address	Land area \
223	4	2.0	10 Eyton Kay Road Hobsonville, Auckland	161.0
309	4	NaN	14 Hea Road Hobsonville, Auckland	214.0
310	4	2.0	12 Hea Road Hobsonville, Auckland	191.0
311	4	NaN	16 Hea Road Hobsonville, Auckland	245.0
449	5	4.0	10 Mantis Lane Hobsonville, Auckland	336.0
460	4	2.0	12 Williams Road Hobsonville, Auckland	450.0
666	3	2.0	17 Kanuka Road Hobsonville, Auckland	259.0

	CV	Latitude	Longitude	SA1	0-19 years	20-29 years \
223	860000	-36.795951	174.655930	7002301	51	36
309	1250000	-36.798371	174.647430	7002267	60	66
310	530000	-36.798371	174.647430	7002267	60	66
311	1100000	-36.798371	174.647430	7002267	60	66
449	545000	-36.801329	174.666149	7002304	33	30
460	1125000	-36.800550	174.645182	7002271	12	12
666	920000	-36.793782	174.660944	7002295	45	24

	30-39 years	40-49 years	50-59 years	60+ years	Suburbs \
223	57	18	6	12	Hobsonville

309	60	24	24	18	Hobsonville
310	60	24	24	18	Hobsonville
311	60	24	24	18	Hobsonville
449	24	15	12	18	Hobsonville
460	9	6	9	75	Hobsonville
666	48	15	6	12	Hobsonville

	C18_CURPop	NZDep2018
223	174	1.0
309	252	2.0
310	252	2.0
311	252	2.0
449	135	4.0
460	129	7.0
666	147	2.0

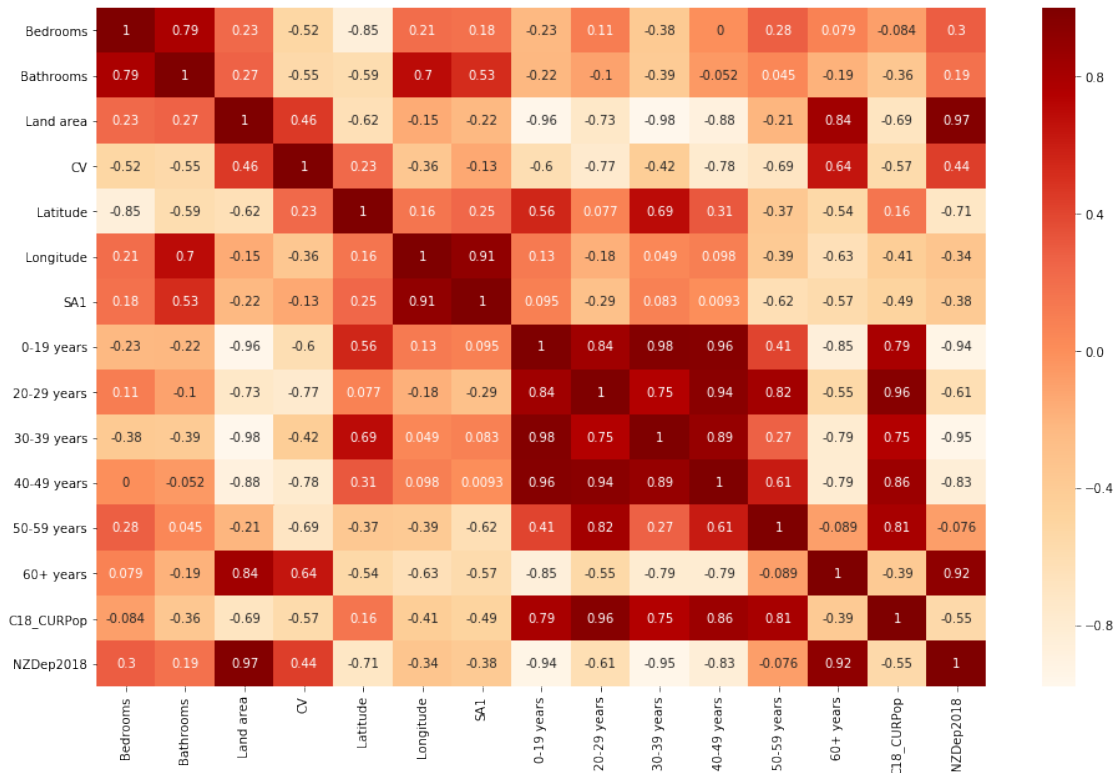
```
[330]: print(house_prices.dtypes)
```

```
Bedrooms          int64
Bathrooms         float64
Address           object
Land area         float64
CV                int64
Latitude          float64
Longitude         float64
SA1              int64
0-19 years        int64
20-29 years       int64
30-39 years       int64
40-49 years       int64
50-59 years       int64
60+ years         int64
Suburbs           object
C18_CURPop        int64
NZDep2018         float64
dtype: object
```

```
[331]: speard_attributes = house_prices.describe()
```

```
[340]: subs_unclean = house_prices[house_prices['Suburbs'] == 'Hobsonville']
subs = subs_unclean.dropna()
```

```
[333]: #using the subs we will try to predict the nan value for bathrooms based on
       →similar houses in the suburbs consider values of 0.5 and higher
ax, fig = plt.subplots(figsize= (16,10))
correlation_matrix = subs.corr()
sns.heatmap(correlation_matrix, annot=True, cmap="OrRd")
plt.show()
```



From the correlation matrix above, we found some features related to CV they now used later for predicting them. These are a houses Bedrooms, Bathrooms and house depression index.

```
[334]: #creating a linear regression model for Bathrooms
feature_names = ['Bedrooms', 'SA1', 'CV', 'Latitude', 'Longitude']
#contains the rows we need to predict to fill in.
input_for_predicting = subs_unclean[sub Unclean['Bathrooms'].isna()]
input_for_predicting
```

```
[334]: Bedrooms Bathrooms Address Land area \
309 4 NaN 14 Hea Road Hobsonville, Auckland 214.0
311 4 NaN 16 Hea Road Hobsonville, Auckland 245.0

CV Latitude Longitude SA1 0-19 years 20-29 years \
309 1250000 -36.798371 174.64743 7002267 60 66
311 1100000 -36.798371 174.64743 7002267 60 66

30-39 years 40-49 years 50-59 years 60+ years Suburbs \
309 60 24 24 18 Hobsonville
311 60 24 24 18 Hobsonville

C18_CURPop NZDep2018
309 252 2.0
311 252 2.0
```



```
[335]: #Machine learning model
#split into training set and test set for after training

from sklearn.model_selection import train_test_split
y = subs[['Bathrooms']]
x = subs[feature_names]
print(x)
print(y)
```

	Bedrooms	SA1	CV	Latitude	Longitude
223	4	7002301	860000	-36.795951	174.655930
310	4	7002267	530000	-36.798371	174.647430
449	5	7002304	545000	-36.801329	174.666149
460	4	7002271	1125000	-36.800550	174.645182
666	3	7002295	920000	-36.793782	174.660944

	Bathrooms
223	2.0
310	2.0
449	4.0
460	2.0
666	2.0

```
[336]: # we have the input and output data ready now we can go about finding linear
→regression model to predict bathroom subur
#import linear regression model for cont data
from sklearn.linear_model import LinearRegression
model = LinearRegression()
```

```
[337]: model.fit(x,y)
```

```
[337]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
[338]: #determine model coefficients
model.coef_
```

```
[338]: array([[ 8.06711279e-01, -3.49804349e-02,  5.87961799e-07,
-3.46078799e+01,  1.26950090e+02]])
```

```
[341]: # from nan row
x_input = input_for_predicting[feature_names]
predicted_bathrooms = model.predict(x_input)
```

```
[342]: predicted_bathrooms = np.round(predicted_bathrooms) # round the float value
```

```
[343]: house_prices.loc[house_prices['Bathrooms'].isnull(), 'Bathrooms'] =
→predicted_bathrooms
house_prices[house_prices['Suburbs'] == 'Hobsonville']
```

```
[343]:
```

	Bedrooms	Bathrooms		Address	Land area \
223	4	2.0	10 Eyton Kay Road	Hobsonville, Auckland	161.0
309	4	2.0	14 Hea Road	Hobsonville, Auckland	214.0
310	4	2.0	12 Hea Road	Hobsonville, Auckland	191.0
311	4	2.0	16 Hea Road	Hobsonville, Auckland	245.0
449	5	4.0	10 Mantis Lane	Hobsonville, Auckland	336.0
460	4	2.0	12 Williams Road	Hobsonville, Auckland	450.0
666	3	2.0	17 Kanuka Road	Hobsonville, Auckland	259.0

	CV	Latitude	Longitude	SA1	0-19 years	20-29 years \
223	860000	-36.795951	174.655930	7002301	51	36
309	1250000	-36.798371	174.647430	7002267	60	66
310	530000	-36.798371	174.647430	7002267	60	66
311	1100000	-36.798371	174.647430	7002267	60	66
449	545000	-36.801329	174.666149	7002304	33	30
460	1125000	-36.800550	174.645182	7002271	12	12
666	920000	-36.793782	174.660944	7002295	45	24

	30-39 years	40-49 years	50-59 years	60+ years	Suburbs \
223	57	18	6	12	Hobsonville
309	60	24	24	18	Hobsonville
310	60	24	24	18	Hobsonville
311	60	24	24	18	Hobsonville
449	24	15	12	18	Hobsonville
460	9	6	9	75	Hobsonville
666	48	15	6	12	Hobsonville

	C18_CURPop	NZDep2018
223	174	1.0
309	252	2.0
310	252	2.0
311	252	2.0
449	135	4.0
460	129	7.0
666	147	2.0

Once again we check if null value exists

```
[344]: last_nan = house_prices[house_prices.isnull().any(axis=1)]
last_nan
```

```
[344]:
```

	Bedrooms	Bathrooms		Address	\
568	1	1.0	14 Te Rangitawhiri Road	Great Barrier Island, ...	

	Land area	CV	Latitude	Longitude	SA1	0-19 years \
568	2141.0	740000	-36.197282	175.416921	7001131	27

	20-29 years	30-39 years	40-49 years	50-59 years	60+ years	Suburbs \
568	6	6	18	39	60	NaN

	C18_CURPop	NZDep2018
568	156	9.0

it's only one entry so just check google maps and checked if in the subset of Surburbs

```
[345]: # Great Barrier Island is the column so we just need to replace that entry.Used
→a google search because it was only one entry and was catgolical
house_prices.loc[house_prices['Suburbs'].isnull(),'Suburbs'] = 'Great Barrier
→Island (Aotea Island)'
```

```
[346]: #complete with the nan completely clean from the dataset
house_prices.isnull().values.any()
```

```
[346]: False
```

Now I will going to look at the spread of each column

```
[347]: def Outlier_range(q1,q3):
        IQR = q3 - q1
        lower = q1 - 1.5 * IQR
        upper = q3 + 1.5 * IQR
        return [lower,upper]
```

```
[348]: #encoding the suburbs as dictionary so that visulation libary's can plot the
→surburbs column
labels = house_prices['Suburbs'].astype('category').cat.categories.tolist()
replace_map_comp = {'Suburbs' : {k: v for k,v in
→zip(labels,list(range(1,len(labels)+1)))}}
replace_map_comp
```

```
[348]: {'Suburbs': {'Albany Heights': 1,
                    'Alfriston': 2,
                    'Army Bay': 3,
                    'Auckland Central': 4,
                    'Avondale': 5,
                    'Beach Haven': 6,
                    'Beachlands': 7,
                    'Belmont': 8,
                    'Birkdale': 9,
                    'Birkenhead': 10,
                    'Blockhouse Bay': 11,
                    'Bombay': 12,
                    'Botany Downs': 13,
                    'Browns Bay': 14,
                    'Buckland': 15,
                    'Bucklands Beach': 16,
                    'Burswood': 17,
                    'Campbells Bay': 18,
                    'Chatswood': 19,
                    'Clarks Beach': 20,
```

'Clendon Park': 21,  
'Clover Park': 22,  
'Cockle Bay': 23,  
'Conifer Grove': 24,  
'Dairy Flat': 25,  
'Dannemora': 26,  
'Drury': 27,  
'East Tamaki': 28,  
'East Tamaki Heights': 29,  
'Eastern Beach': 30,  
'Eden Terrace': 31,  
'Ellerslie': 32,  
'Epsom': 33,  
'Farm Cove': 34,  
'Favona': 35,  
'Flat Bush': 36,  
'Forrest Hill': 37,  
'Freemans Bay': 38,  
'Glen Eden': 39,  
'Glen Innes': 40,  
'Glendene': 41,  
'Glendowie': 42,  
'Glenfield': 43,  
'Golflands': 44,  
'Goodwood Heights': 45,  
'Great Barrier Island (Aotea Island)': 46,  
'Green Bay': 47,  
'Greenlane': 48,  
'Grey Lynn': 49,  
'Gulf Harbour': 50,  
'Half Moon Bay': 51,  
'Hatfields Beach': 52,  
'Helensville': 53,  
'Henderson': 54,  
'Herne Bay': 55,  
'Highland Park': 56,  
'Hillcrest': 57,  
'Hillsborough': 58,  
'Hobsonville': 59,  
'Howick': 60,  
'Huapai': 61,  
'Huia': 62,  
'Hunua': 63,  
'Kaipara Flats': 64,  
'Karaka': 65,  
'Kawakawa Bay': 66,  
'Kelston': 67,

'Kingsland': 68,  
'Kohimarama': 69,  
'Kumeu': 70,  
'Laingholm': 71,  
'Leigh': 72,  
'Long Bay': 73,  
'Lynfield': 74,  
'Mangere': 75,  
'Mangere Bridge': 76,  
'Mangere East': 77,  
'Manukau': 78,  
'Manurewa': 79,  
'Manurewa East': 80,  
'Maraetai': 81,  
'Massey': 82,  
'Matakatia': 83,  
'Mauku': 84,  
'Meadowbank': 85,  
'Mellons Bay': 86,  
'Milford': 87,  
'Mission Bay': 88,  
'Morningside': 89,  
'Mount Albert': 90,  
'Mount Eden': 91,  
'Mount Roskill': 92,  
'Mount Wellington': 93,  
'Murrays Bay': 94,  
'Narrow Neck': 95,  
'New Lynn': 96,  
'New Windsor': 97,  
'Northcote': 98,  
'Okura': 99,  
'Omiha': 100,  
'One Tree Hill': 101,  
'Onehunga': 102,  
'Oneroa': 103,  
'Onetangi': 104,  
'Opaheke': 105,  
'Orakei': 106,  
'Orewa': 107,  
'Ostend': 108,  
'Otahuhu': 109,  
'Otara': 110,  
'Oteha': 111,  
'Paerata': 112,  
'Pahurehure': 113,  
'Pakuranga': 114,

'Pakuranga Heights': 115,  
'Palm Beach': 116,  
'Panmure': 117,  
'Papakura': 118,  
'Papatoetoe': 119,  
'Parakai': 120,  
'Paremoremo': 121,  
'Parnell': 122,  
'Patumahoe': 123,  
'Penrose': 124,  
'Pinehill': 125,  
'Pohuehue': 126,  
'Point Chevalier': 127,  
'Point England': 128,  
'Pokeno': 129,  
'Ponsonby': 130,  
'Pukekohe': 131,  
'Rakino Island': 132,  
'Ramarama': 133,  
'Randwick Park': 134,  
'Ranui': 135,  
'Red Beach': 136,  
'Red Hill': 137,  
'Redvale': 138,  
'Remuera': 139,  
'Rosehill': 140,  
'Rothesay Bay': 141,  
'Royal Oak': 142,  
'Saint Johns': 143,  
'Saint Marys Bay': 144,  
'Sandringham': 145,  
'Schnapper Rock': 146,  
'Silverdale': 147,  
'Snells Beach': 148,  
'Somerville': 149,  
'South Head': 150,  
'St Heliers': 151,  
'Stanmore Bay': 152,  
'Stonefields': 153,  
'Sunnyhills': 154,  
'Sunnyvale': 155,  
'Surfdale': 156,  
'Swanson': 157,  
'Takanini': 158,  
'Te Atatu Peninsula': 159,  
'Te Atatu South': 160,  
'The Gardens': 161,

```

'Three Kings': 162,
'Tindalls Beach': 163,
'Titirangi': 164,
'Torbay': 165,
'Totara Heights': 166,
'Totara Park': 167,
'Totara Vale': 168,
'Tuakau': 169,
'Unsworth Heights': 170,
'Wade Heads': 171,
'Wai O Taiki Bay': 172,
'Waiheke Island': 173,
'Waimauku': 174,
'Wainui': 175,
'Waitakere': 176,
'Waitoki': 177,
'Waiuku': 178,
'Warkworth': 179,
'Waterview': 180,
'Wattle Downs': 181,
'Wellsford': 182,
'Wesley': 183,
'West Harbour': 184,
'Westmere': 185,
'Weymouth': 186,
'Whenuapai': 187,
'Windsor Park': 188,
'Wiri': 189}}

```

```
[349]: #making a copy to modify the suburbs data directly
```

```
house_prices_replace = house_prices.copy()
```

```
[350]: #replace the numeric values in dataframe for suburbs
```

```
house_prices_replace.replace(replace_map_comp, inplace=True)
```

```
[351]: house_prices_replace.dtypes
```

```

[351]: Bedrooms          int64
      Bathrooms         float64
      Address           object
      Land area         float64
      CV                int64
      Latitude          float64
      Longitude         float64
      SA1               int64
      0-19 years        int64
      20-29 years       int64
      30-39 years       int64
      40-49 years       int64

```

```

50-59 years      int64
60+ years        int64
Suburbs          int64
C18_CURPop       int64
NZDep2018        float64
dtype: object

```

```

[352]: #dropping the address column because suburbs can act as more generic input for
        →our model
house_prices_replace.drop(['Address'],axis=1).head()

```

```

[352]:   Bedrooms  Bathrooms  Land area      CV  Latitude  Longitude      SA1  \
0         5         3.0      714.0  960000 -37.012920  174.904069  7009770
1         5         3.0      564.0 1250000 -37.063672  174.922912  7009991
2         6         4.0      626.0 1250000 -37.063580  174.924044  7009991
3         2         1.0       65.0  740000 -36.912996  174.787425  7007871
4         3         1.0      601.0  630000 -36.979037  174.892612  7008902

      0-19 years  20-29 years  30-39 years  40-49 years  50-59 years  60+ years  \
0             48           27           24           21           24           21
1             42           18           12           21           15           30
2             42           18           12           21           15           30
3             42            6           21           21           12           15
4             93           27           33           30           21           33

      Suburbs  C18_CURPop  NZDep2018
0          79         174         6.0
1          65         129         1.0
2          65         129         1.0
3         102         120         2.0
4          22         231         9.0

```

```

[353]: stats_house_prices_replace= house_prices_replace.describe()
stats_house_prices_replace
house_prices_replace.shape
house_prices_replace.describe()

```

```

[353]:   count      Bedrooms      Bathrooms      Land area      CV      Latitude  \
count  1051.000000  1051.000000  1051.000000  1.051000e+03  1051.000000
mean     3.777355     2.073264     856.989534  1.387521e+06  -36.893715
std     1.169412     0.992044    1588.156219  1.182939e+06     0.130100
min     1.000000     1.000000     40.000000  2.700000e+05  -37.265021
25%     3.000000     1.000000     321.000000  7.800000e+05  -36.950565
50%     4.000000     2.000000     571.000000  1.080000e+06  -36.893132
75%     4.000000     3.000000     825.000000  1.600000e+06  -36.855789
max     17.000000     8.000000    22240.000000  1.800000e+07  -36.177655

      Longitude      SA1  0-19 years  20-29 years  30-39 years  \
count  1051.000000  1.051000e+03  1051.000000  1051.000000  1051.000000

```



mean	174.799325	7.006319e+06	47.549001	28.963844	27.042816
std	0.119538	2.591262e+03	24.692205	21.037441	17.975408
min	174.317078	7.001130e+06	0.000000	0.000000	0.000000
25%	174.720779	7.004416e+06	33.000000	15.000000	15.000000
50%	174.798575	7.006325e+06	45.000000	24.000000	24.000000
75%	174.880944	7.008384e+06	57.000000	36.000000	33.000000
max	175.492424	7.011028e+06	201.000000	270.000000	177.000000

	40-49 years	50-59 years	60+ years	Suburbs	C18_CURPop \
count	1051.000000	1051.000000	1051.000000	1051.000000	1051.000000
mean	24.125595	22.615604	29.360609	94.834443	179.914367
std	10.942770	10.210578	21.805031	50.446475	71.059280
min	0.000000	0.000000	0.000000	1.000000	3.000000
25%	18.000000	15.000000	18.000000	54.000000	138.000000
50%	24.000000	21.000000	27.000000	93.000000	174.000000
75%	30.000000	27.000000	36.000000	139.000000	210.000000
max	114.000000	90.000000	483.000000	189.000000	789.000000

	NZDep2018
count	1051.000000
mean	5.063749
std	2.913471
min	1.000000
25%	2.000000
50%	5.000000
75%	8.000000
max	10.000000

Initial Data Exploration The initial exploration of the data began with the collecting of current population value through an API then merging the existing data frame and combining with depreciation index of houses. Then NAN values in the dataset were replaced with estimates. Finally along with converting a suburbs column into category variable. The descriptive statistics summary was used to calculate outliers. This then reduced the number of rows to 1051—shown bellow with an interquartile estimate.

```
[354]: #removing outliers from Land Area CV these should also remove the houses rows
        →larger amounts bedrooms
q1 = stats_house_prices_replace['Land area'].loc['25%']
q3 = stats_house_prices_replace['Land area'].loc['75%']
outrang = Outlier_range(q1,q3)
print(outrang)
house_prices_replace = house_prices_replace[house_prices_replace['Land area'] <
        →outrang[1]] # check if outside the outer range
```

```
[-435.0, 1581.0]
```

```
[355]: house_prices_replace.describe()
```

```
[355]:
```

	Bedrooms	Bathrooms	Land area	CV	Latitude	\
count	981.000000	981.000000	981.000000	9.810000e+02	981.000000	
mean	3.733945	2.047910	560.669725	1.326148e+06	-36.896886	
std	1.053449	0.955515	334.873663	1.003560e+06	0.120351	
min	1.000000	1.000000	40.000000	2.700000e+05	-37.265021	
25%	3.000000	1.000000	304.000000	7.700000e+05	-36.949429	
50%	4.000000	2.000000	536.000000	1.050000e+06	-36.894835	
75%	4.000000	3.000000	763.000000	1.550000e+06	-36.857091	
max	8.000000	7.000000	1573.000000	1.800000e+07	-36.298627	

	Longitude	SA1	0-19 years	20-29 years	30-39 years	\
count	981.000000	9.810000e+02	981.000000	981.000000	981.000000	
mean	174.797593	7.006361e+06	47.828746	29.483180	27.623853	
std	0.107924	2.538859e+03	25.147103	21.451249	18.260256	
min	174.433162	7.001158e+06	0.000000	0.000000	0.000000	
25%	174.723850	7.004537e+06	33.000000	18.000000	15.000000	
50%	174.797892	7.006334e+06	45.000000	27.000000	24.000000	
75%	174.879070	7.008379e+06	57.000000	36.000000	36.000000	
max	175.187565	7.011028e+06	201.000000	270.000000	177.000000	

	40-49 years	50-59 years	60+ years	Suburbs	C18_CURPop	\
count	981.000000	981.000000	981.000000	981.000000	981.000000	
mean	24.024465	22.244648	28.993884	94.107034	180.458716	
std	11.017677	10.016175	22.196662	50.093645	72.527219	
min	0.000000	0.000000	0.000000	3.000000	3.000000	
25%	18.000000	15.000000	18.000000	54.000000	138.000000	
50%	24.000000	21.000000	27.000000	92.000000	174.000000	
75%	30.000000	27.000000	36.000000	139.000000	210.000000	
max	114.000000	90.000000	483.000000	189.000000	789.000000	

	NZDep2018
count	981.000000
mean	5.160041
std	2.897549
min	1.000000
25%	3.000000
50%	5.000000
75%	8.000000
max	10.000000

```
[356]: house_prices_replace.shape
# removed some of rows which outliers
```

```
[356]: (981, 17)
```

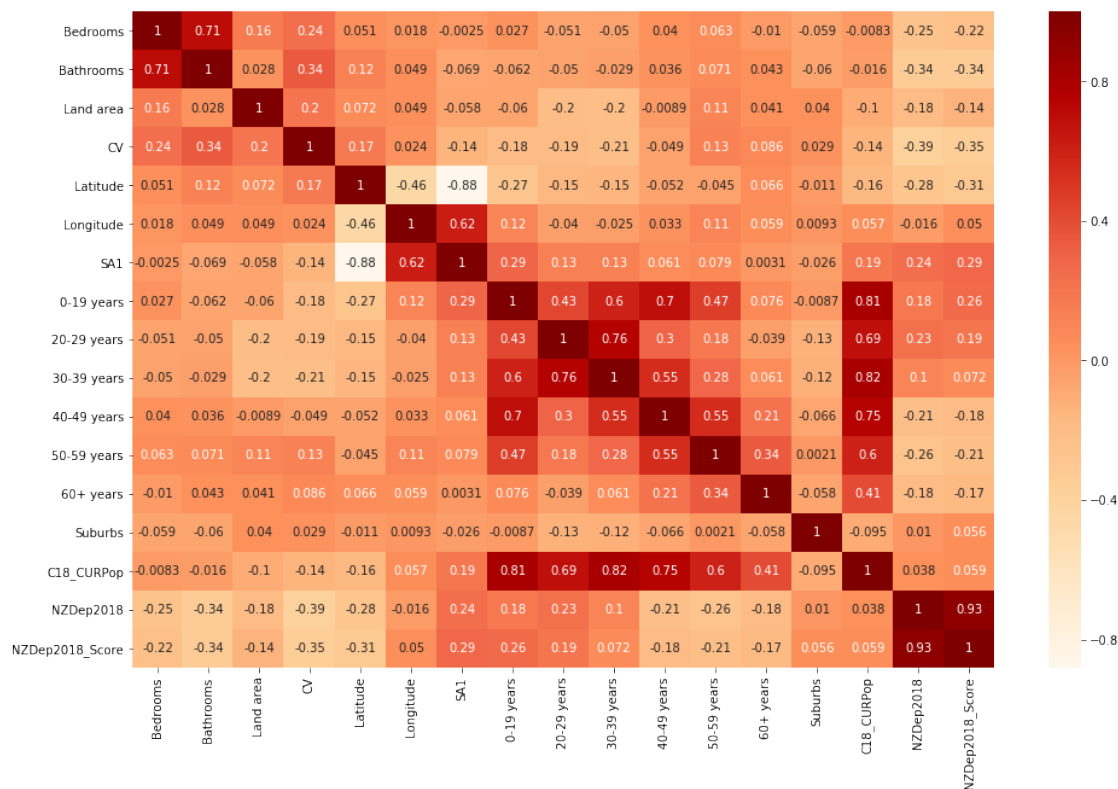
Their exists a continuous version of NZDEP index otago dataset so adding it to house\_prices dataset see if releationship could be found.

```
[358]: def continuous_extract_depreciation(df_dep,sa1):
        # get row where sa1 has the same value
        row = df_dep.loc[df_dep['SA12018_code'] == sa1]
        return float(row['NZDep2018_Score'])
```

```
[359]: house_prices_replace['NZDep2018_Score'] = house_prices_replace.apply(lambda row :
        → continuous_extract_depreciation(dep_indexs,row['SA1']),axis = 1)
```

Correlation and Relationships Numeric Relationships The correlation between the numeric columns was calculated and observed in the below correlation plot. (The right colour bar indicated the correlation values. For example, the dark red means correlation value is 1, and pale yellow means correlation value is negative 1.) The graph shows that Bedrooms, Bathrooms and NZDep2018 correlate with CV.

```
[360]: # check the collection matrix now added the continous depression
ax, fig = plt.subplots(figsize= (16,10))
correlation_matrix = house_prices_replace.corr()
sns.heatmap(correlation_matrix, annot=True, cmap="OrRd")
plt.show()
```



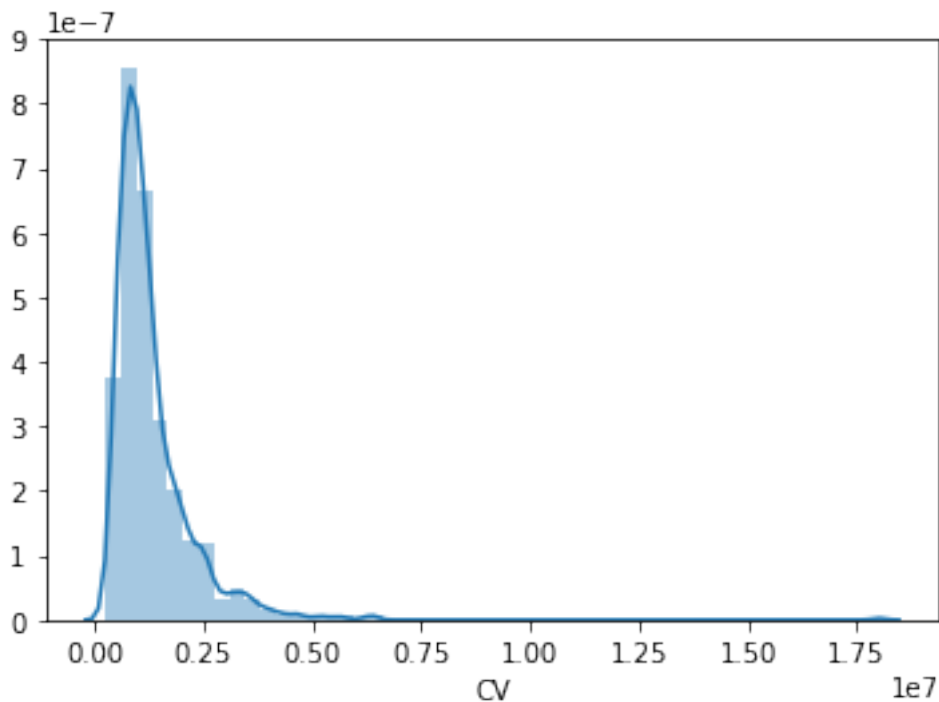
Bellow using the correlation matrix for CV visulations for, releated variables were created and analyzed.

```
[361]: sns.distplot(house_prices_replace['CV'])
```

```
/home/nbuser/anaconda3_501/lib/python3.6/site-  
packages/scipy/stats/stats.py:1713: FutureWarning: Using a non-tuple sequence  
for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of  
`arr[seq]`. In the future this will be interpreted as an array index,  
`arr[np.array(seq)]`, which will result either in an error or a different  
result.
```

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

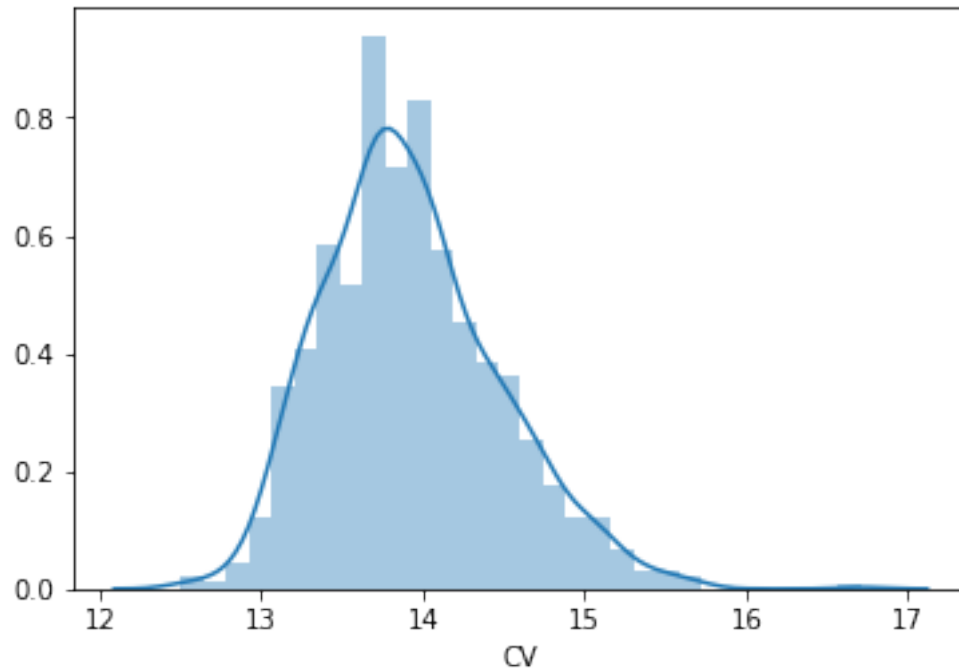
[361]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f3c7e3037b8>



Price of the house is right-skewed, so log transformation was applied resulting in more normally distributed data

```
[362]: sns.distplot(np.log(house_prices_replace['CV']))
```

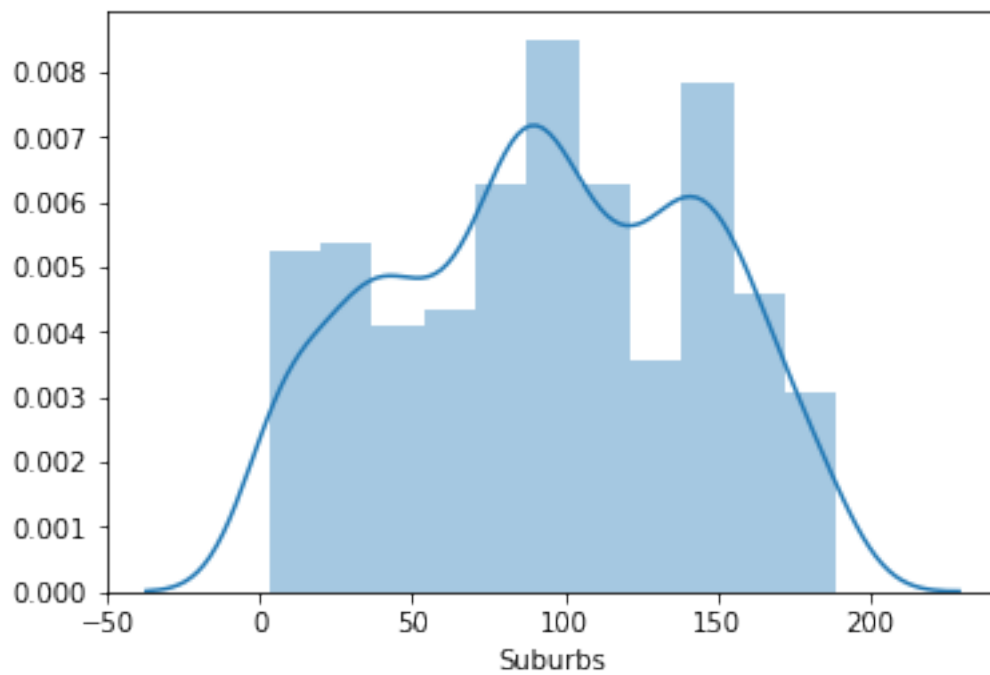
[362]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f3c7eb02470>



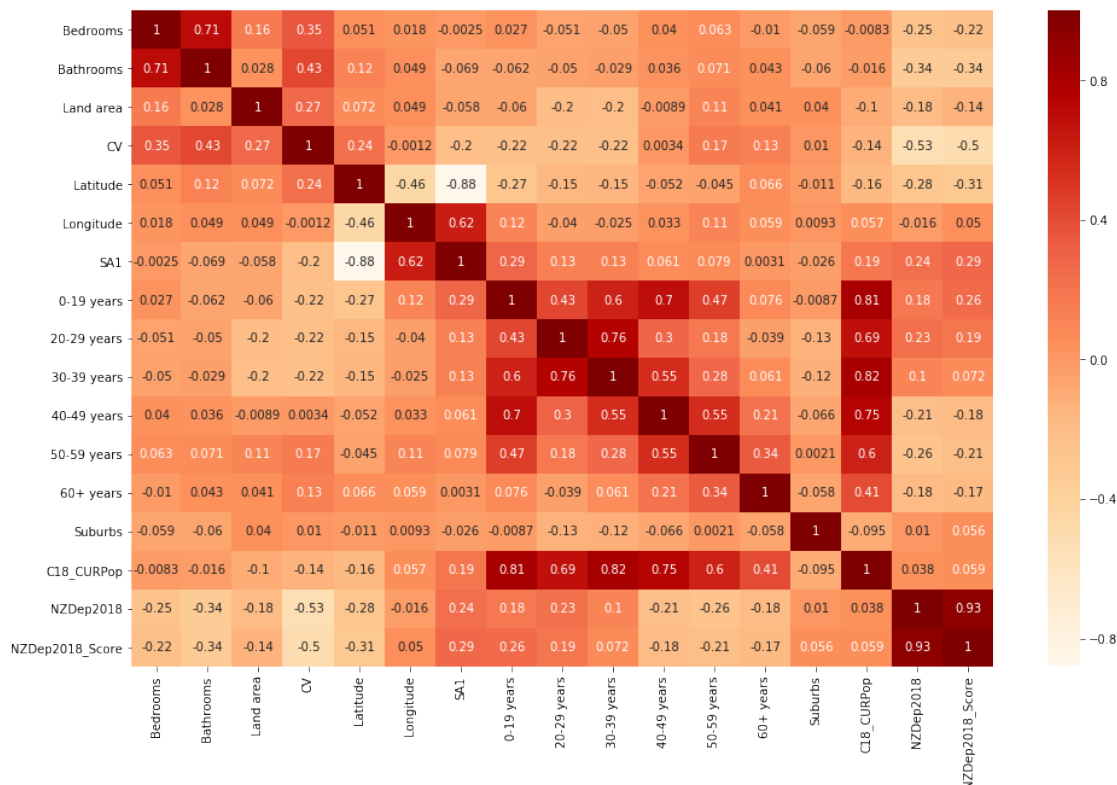
```
[363]: # now the price is in log form
house_prices_replace['CV'] = house_prices_replace['CV'].apply(np.log)
```

```
[364]: sns.distplot(house_prices_replace['Suburbs'])
```

```
[364]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3c7cbae400>
```



```
[365]: # check the collction matrix now added the continous depression
ax, fig = plt.subplots(figsize= (16,10))
correlation_matrix = house_prices_replace.corr()
sns.heatmap(correlation_matrix, annot=True, cmap="OrRd")
plt.show()
```



The above correlation matrix indicates CV is related negatively with NZDep index value and positively associated with Bathroom area, hinting at the higher number of bathrooms in the house, the higher it estimated cost. So we will now develop the linear regression model for CV based on NZDep index bedrooms, Bathrooms. Note that the output of CV is log form, so when using the linear model for predications, it required to exp the output to get house price.

#### Analysis

In this analysis, Linear regression was tested. These algorithms were trained with 40% of the data. Testing the model with the remaining 60% of the data yielded the following results:

```
[366]: #define the inputs x contains NZDep2018 data along with Bathroom data
#define y out data being log of CV
x = house_prices_replace[['Bedrooms', 'Bathrooms', 'NZDep2018']]
y = house_prices_replace['CV']
#dealing with 981 data points per column
len(x)
```

[366]: 981

```
[367]: #splitting the data into training and predicting batchs
train_x , test_x, train_y, test_y = train_test_split(x,y,test_size = 0.4,
↳random_state=42)
```

```
[368]: #define a new regression linear model
model_CV = LinearRegression()
```

```
[369]: model_CV.fit(train_x,train_y)
```

```
[369]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
[370]: #determine model coefficients
model_CV.coef_
```

```
[370]: array([ 0.02176466,  0.15031556, -0.08151677])
```

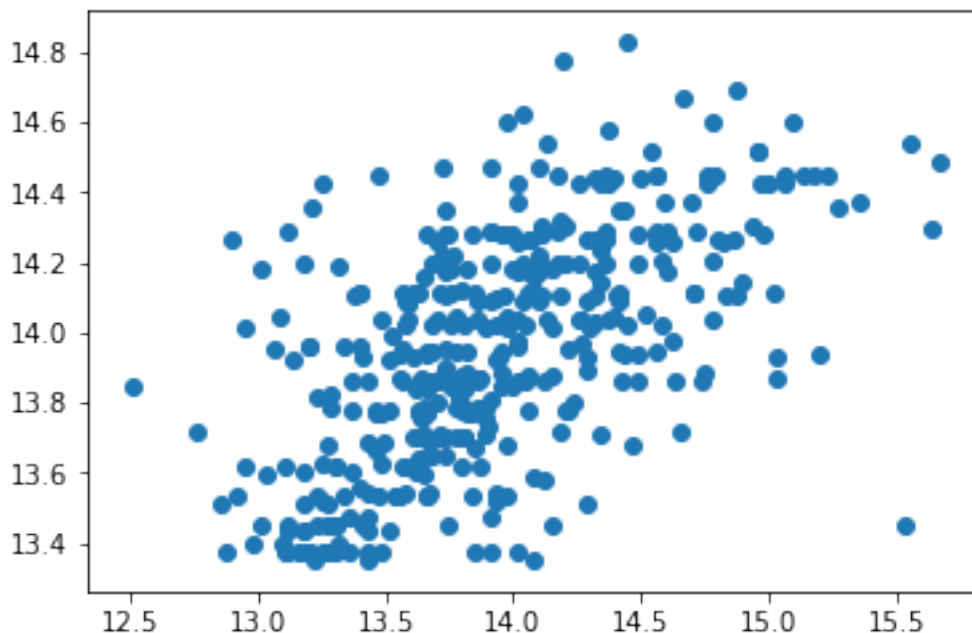
```
[372]: predicted = model_CV.predict(test_x)
```

Analysis The decision for the model chosen, process and results

Conclusion This analysis has shown that the (house price) in log scale prediction can not be confidently predicted from it's a number of Bedrooms, Bathrooms and property depreciation index. The accuracy rate is 33% which bellow 50 % and scatter plot has revealed house price doesn't follow a linear trend there might need to be some additional work need with cleaning the input into the model or transformations need on the inputs.

```
[373]: plt.scatter(test_y,predicted)
```

```
[373]: <matplotlib.collections.PathCollection at 0x7f3c6abc6898>
```



```
[374]: model_CV.score(test_x, test_y)
```

```
[374]: 0.33913597501995824
```

The model above give a score of 0.34. which similar to correlated values in correlation matrix. However from looking at the scatter plot it looks like data is very sparse it indicates it more cleaning would need to be done with input into the model.